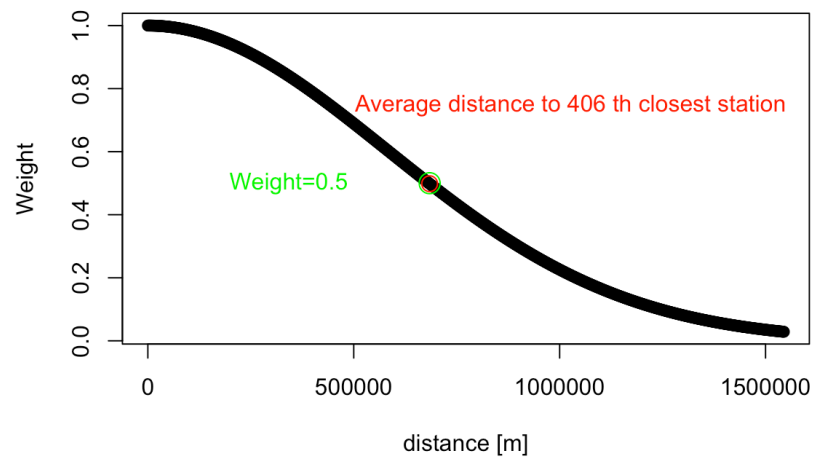
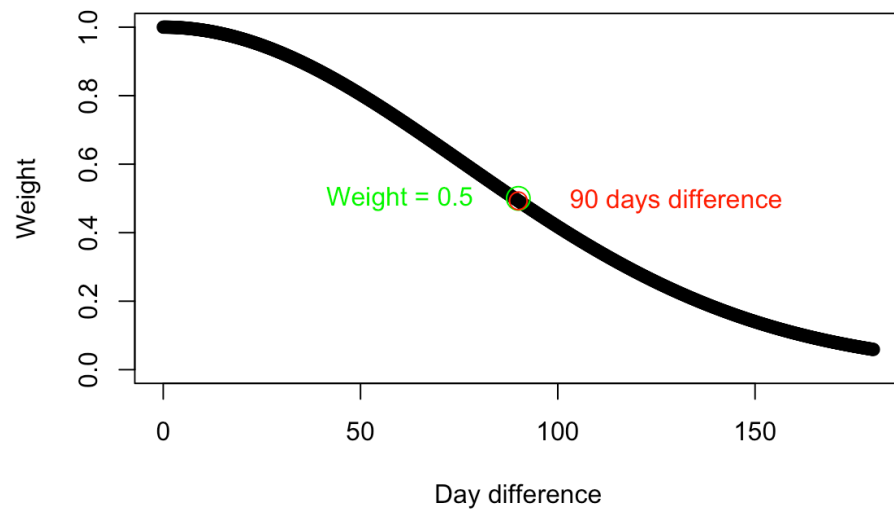
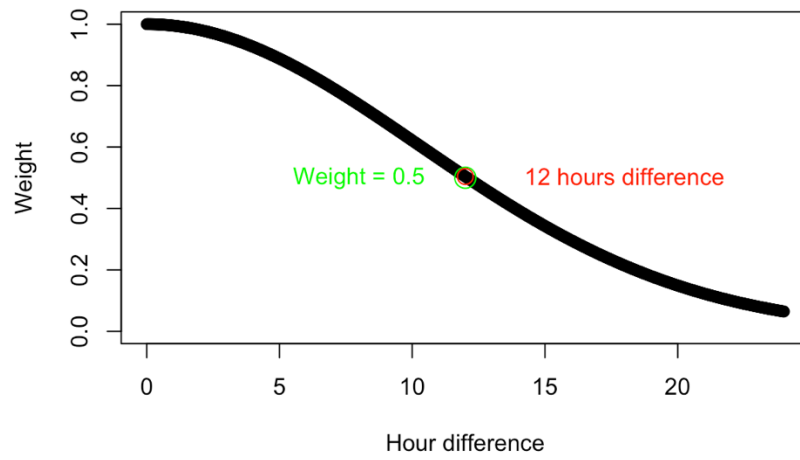


Assignment 1: Kernel Methods



The width of kernels were chosen so that the weight had decreased to 0.5 at a certain distance I figured was reasonable.

Most weighted predictions at 08:00:00 with summed kernels:

distance	date	time
"64209.3475626027"	"2012-12-28"	"08:00:00"
distance	date	time
"67647.8518002752"	"1998-12-26"	"09:00:00"
distance	date	time
"64428.0670248181"	"1994-12-20"	"09:00:00"

Most weighted predictions at 08:00:00 with multiplied kernels:

distance	date	time
"64209.3475626027"	"2012-12-28"	"08:00:00"
distance	date	time
"67647.8518002752"	"1998-12-26"	"09:00:00"
distance	date	time
"64428.0670248181"	"1994-12-20"	"09:00:00"

Most weighted predictions at 14:00:00 with summed kernels:

distance	date	time
"55443.5735022029"	"2011-12-31"	"14:00:00"
distance	date	time
"55443.5735022029"	"2001-12-24"	"15:00:00"
distance	date	time
"64209.3475626027"	"2010-12-22"	"15:00:00"

Most weighted predictions at 14:00:00 with multiplied kernels:

distance	date	time
"55443.5735022029"	"2011-12-31"	"14:00:00"
distance	date	time
"55443.5735022029"	"2001-12-24"	"15:00:00"
distance	date	time
"64209.3475626027"	"2010-12-22"	"15:00:00"

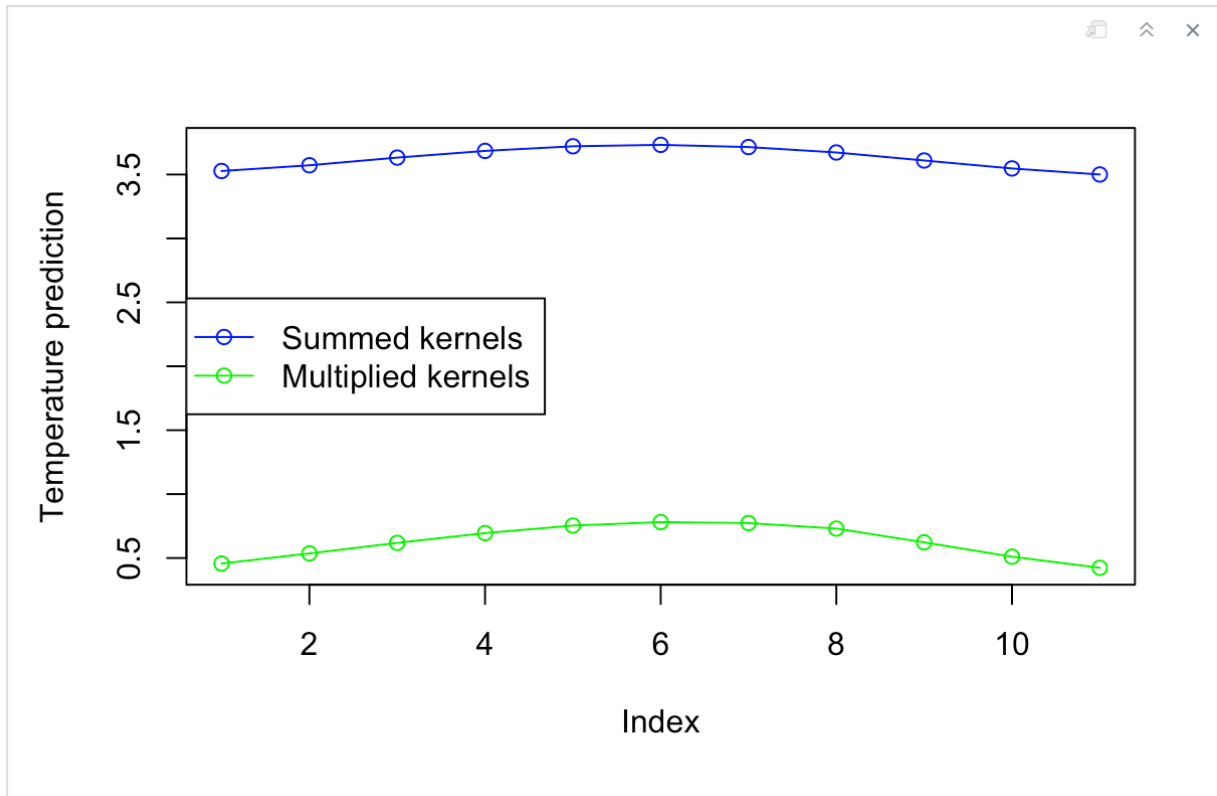
Most weighted predictions at 20:00:00 with summed kernels:

distance	date	time
"67647.8518002752"	"2007-12-25"	"21:00:00"
distance	date	time
"67647.8518002752"	"1973-12-16"	"21:00:00"
distance	date	time
"55443.5735022029"	"2009-12-23"	"18:00:00"

Most weighted predictions at 20:00:00 with multiplied kernels:

distance	date	time
----------	------	------

"67647.8518002752"	"2007-12-25"	"21:00:00"
distance	date	time
"67647.8518002752"	"1973-12-16"	"21:00:00"
distance	date	time
"55443.5735022029"	"2009-12-23"	"18:00:00"



The contribution from each temperature measurement with the summed kernel gets a higher average weight, compared to the multiplied. This is because if any of the kernels generate weight near zero or equal to zero, the entire weight for that measurement becomes zero with the multiplied kernel

Assignment 3: Neural Networks

Question Assignment 3: Train neural networks to learn the trigonometric sine function. Perform early stop of the gradient descent using the threshold argument. Consider thresholds $i/10000$, where $i = 1 \dots 10$. Initialize the weights of the neural networks to random values in the interval $[-1,1]$. Use a hidden layer of 10 units. Choose the most appropriate value for the threshold. Motivate your choice.

Answer Assignment 3: To answer the question, the library neuralnet was used (as suggested in assignment). To evaluate the different neuralnets, the least square of model predictions was calculated for each model, and the model that is most appropriate is the model with the least square, for the validation set. The least square for each model is plotted in the following plot:

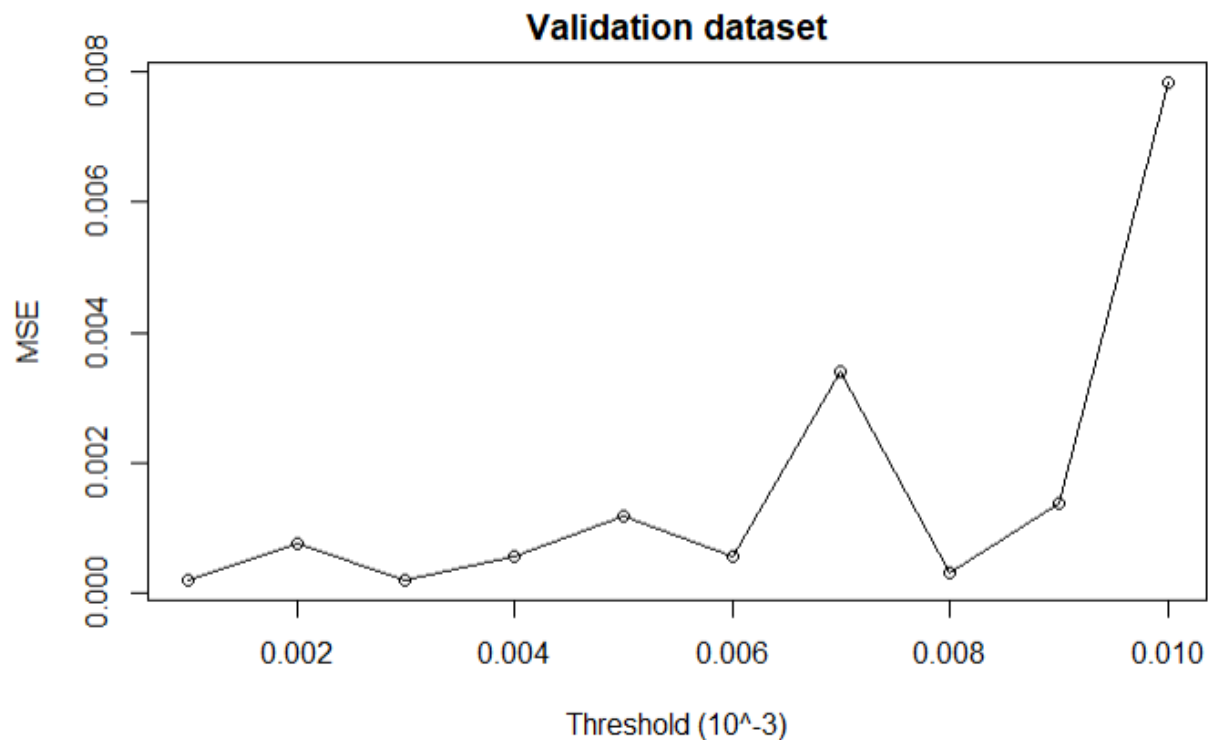


Figure 1, the squared error of model predictions vs. true values for validation set, for different thresholds

As we can see, the model performance differs for different thresholds. The smallest squared error is identified by:

`which.min(SSE)`

Where SSE is the vector containing the squared errors, indexed after i used in threshold. The result is that index 1 has the lowest SSE, with $SSE[1] = 0.00521$ and threshold = $1/1000$. The resulting neural network has the following structure:


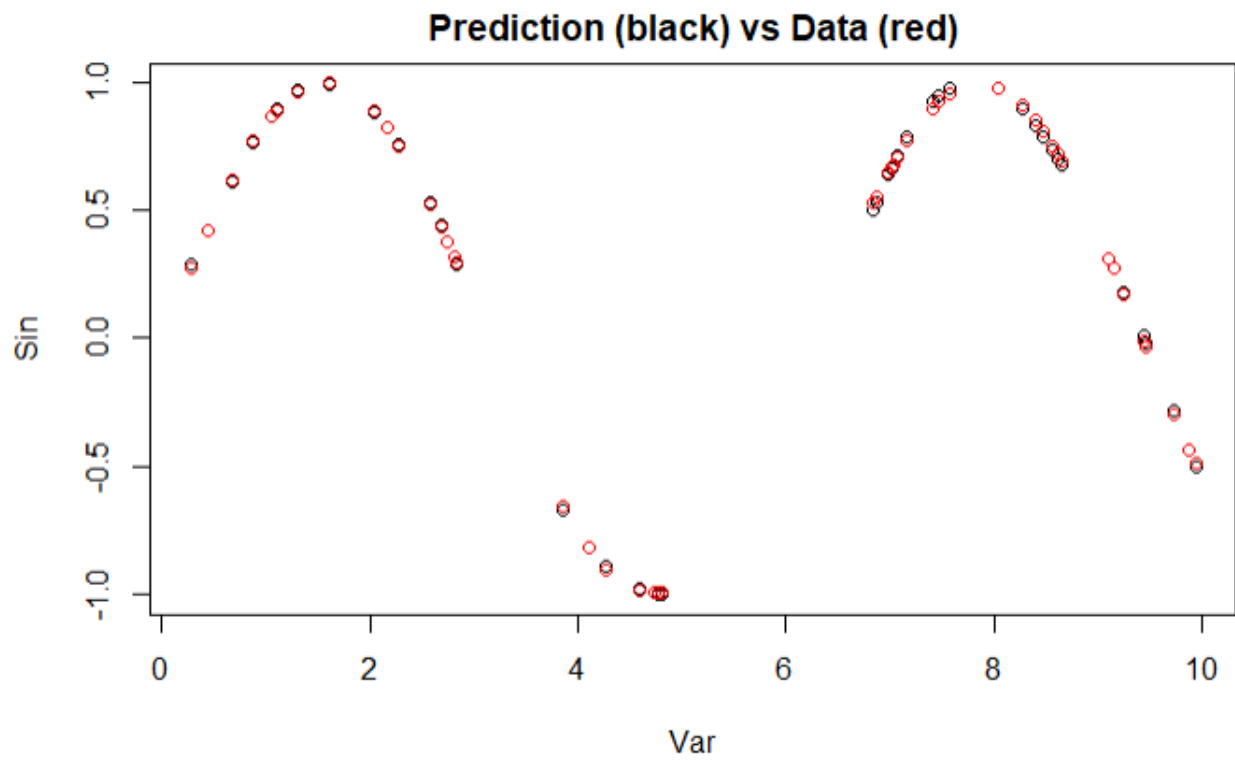


Figure 2, the structure of the best neural network

The resulting function has the following plot:



Appendix 1: code for assignment 3

```
#####
```

```
## 1.1 ##
```

```
#####
```

DETERMING PARAMETERS

```
set.seed(1234567890)
```

```
library(geosphere)
```

```
gaussian_kernel <- function(distance, h){  
  return (exp(-(distance/h)^2))  
}
```

```
#https://en.wikipedia.org/wiki/Geographical\_center\_of\_Sweden
```

```
lat <- 62.173276
```

```
long <- 14.942265
```

```
stations <- read.csv(file="stations.csv", fileEncoding="latin1", stringsAsFactors=FALSE)
```

```
stations <- cbind(stations,distHaversine(p1=stations[,c("latitude","longitude")],p2= c(lat,long)))
```

```
colnames(stations)[ncol(stations)] <- "distance"
```

```
temps <- read.csv(file="temps50k.csv", stringsAsFactors=FALSE)
```

```
st <- merge(stations,temps,by="station_number")
```

```
edge_stations <- vector(length=4)
```

```
edge_stations[1] <- stations[which.max(stations$latitude),"station_number"]
```

```
edge_stations[2] <- stations[which.min(stations$latitude),"station_number"]
```

```

edge_stations[3] <- stations[which.max(stations$longitude),"station_number"]
edge_stations[4] <- stations[which.min(stations$longitude),"station_number"]

threshold_station <- round(nrow(stations)/2)
max_distance <- 0

threshold_distances <- vector(length = length(edge_stations))

for (i in 1:length(edge_stations)){
  distances <- vector(length = nrow(stations))
  for (j in 1:nrow(stations)){
    distances[j] <-
distHaversine(stations[which(stations$station_number==edge_stations[i]),c("longitude","latitude")],
stations[j,c("longitude","latitude")])
  }
  distances <- sort(distances,decreasing = FALSE)
  threshold_distances[i] <- distances[threshold_station]
  max_distance <- round(ifelse(max_distance<max(distances),max(distances), max_distance))
}
threshold_distance <- mean(threshold_distances)

h_distance <- 820000

plot(x = seq(0,max_distance,length.out = 1000), y = gaussian_kernel(distance =
seq(0,max_distance,length.out = 1000), h = h_distance), ylab = "Weight", xlab = "distance [m]")
points(x=threshold_distance, y = gaussian_kernel(distance = threshold_distance,h_distance), col="red",
type = "p", cex=1.5)
text(x=threshold_distance*1.5, y = gaussian_kernel(distance = threshold_distance,h_distance)*1.5,
col="red",labels= paste("Average distance to", threshold_station,"th closest station"))
points(x=threshold_distance, y = 0.5, col="green", type = "p", cex=2)

```

```
text(x=threshold_distance/2, y = 0.5, col="green", labels = "Weight=0.5")
```

```
h_date <- 107
```

```
plot(x=seq(0,180,length.out = 1000), y = gaussian_kernel(seq(0,180,length.out = 1000), h_date), ylab =  
"Weight", xlab = "Day difference", xlim=c(0,180), ylim=c(0,1))
```

```
points(x=90, y = 0.5, col="green", type = "p", cex=2)
```

```
text(x=60, y = 0.5, col="green", labels = "Weight = 0.5")
```

```
points(x=90, y = gaussian_kernel(90, h_date), col="red", type = "p", cex=1.5)
```

```
text(x=130, y = gaussian_kernel(90, h_date), col="red", labels = "90 days difference")
```

```
h_time <- 14.5
```

```
plot(x=seq(0,24,length.out = 1000), y = gaussian_kernel(seq(0,24,length.out = 1000), h_time), ylab =  
"Weight", xlab = "Hour difference", xlim=c(0,24), ylim=c(0,1))
```

```
points(x=12, y = 0.5, col="green", type = "p", cex=2)
```

```
text(x=8, y = 0.5, col="green", labels = "Weight = 0.5")
```

```
points(x=12, y = gaussian_kernel(12, h_time), col="red", type = "p", cex=1.5)
```

```
text(x=18, y = gaussian_kernel(12, h_time), col="red", labels = "12 hours difference")
```

Predicting temperatures

```
date <- "2015-12-24" # The date to predict (up to the students)
```

```
times <- c("04:00:00", "06:00:00",  
"08:00:00", "10:00:00", "12:00:00", "14:00:00", "16:00:00", "18:00:00", "20:00:00", "22:00:00", "24:00:00")
```

```
temp1 <- vector(length=length(times))
```

```
temp2 <- vector(length=length(times))
```



```
# Students' code here
```

```
st <- st[which(st$date<=date),]
```

```
for (t in 1:length(times))
```

```
{  
  exclude <- which((st$date==date)&(st$time>=times[t]))  
  if (length(exclude)>0) {  
    data <- st[-exclude,]  
  } else {  
    data <- st  
  }  
}
```

```
rownames(data) <- 1:nrow(data)
```

```
my_matrix <- matrix(NA,nrow=nrow(data), ncol=4)
```

```
colnames(my_matrix) <- c("temperature", "k_sum", "k_prod", "index")
```

```
for (s in 1:nrow(data)){
```

```
  k1 <- gaussian_kernel(data[s,"distance"], h_distance)
```

```
  days_diff <- as.numeric(as.Date(date)-as.Date(data[s,"date"]))%%365.25
```

```
  days_diff <- min(365-days_diff, days_diff)
```

```
  k2 <- gaussian_kernel(days_diff, h_date)
```

```
  hour_diff <- as.numeric(strptime(times[t],"%H:%M:%S")-  
strptime(data[s,"time"],"%H:%M:%S"))%%24
```

```
  hour_diff <- min(24-hour_diff, hour_diff)
```

```

k3 <- gaussian_kernel(hour_diff,h_time)
my_matrix[s,] <- c(data[s,"air_temperature"], k1+k2+k3, k1*k2*k3, s)
}

```

```

my_matrix[, "k_sum"] <- my_matrix[, "k_sum"]/sum(my_matrix[, "k_sum"])
my_matrix[, "k_prod"] <- my_matrix[, "k_prod"]/sum(my_matrix[, "k_prod"])

```

```

predicton <- my_matrix[, "temperature"]*my_matrix[, "k_sum"]
temp1[t] <- sum(predicton)
predicton <- my_matrix[, "temperature"]*my_matrix[, "k_prod"]
temp2[t] <- sum(predicton)

```

```

my_matrix <- my_matrix[order(my_matrix[, "k_sum"], decreasing = TRUE),]

```

```

if (t%%3 == 0) {

```

```

  cat("\n\nMost weighted predictions at ", times[t], "with summed kernels:\n")

```

```

  for (i in 1:3)
  {
    print(unlist(data[my_matrix[i, "index"], c("distance", "date", "time")]))
  }

```

```

my_matrix <- my_matrix[order(my_matrix[, "k_prod"], decreasing = TRUE),]

```

```

cat("\n\nMost weighted predictions at ", times[t], "with multiplied kernels:\n")

```

```

for (i in 1:3)
{

```

```

        print(unlist(data[my_matrix[i,"index"],c("distance","date", "time")]))
    }
}
}

```

Plot

```

ylim = c(min(min(temp1),min(temp2)), max(max(temp1),max(temp2)))
plot(temp1, type="o", ylim=ylim, col = "blue", ylab="Temperature prediction")
legend("left", legend=c("Summed kernels", "Multiplied kernels"),col=c("blue", "green"), lty = c(1,1), pch
= c(1,1))
lines(temp2, type="o", col = "green")

```

```

``{r}

```

Assignment 3

```

library("neuralnet")

```

```

set.seed(1234567890)

```

```

Var = runif(50, 0, 10)

```

```

data = data.frame(Var, Sin=sin(Var))

```

```

train = data[1:25,]

```

```

validation = data[26:50,]

```

```

# Random initialization of the weights in the interval [-1, 1]

```

```

winit = runif(10,-1,1)

```

```

mse = function(prediction, observation)

```

```

{

```

```
    return (mean((observation - prediction)^2))
}
```

```
n = 10
```

```
mse_val = numeric()
```

```
mse_train = numeric()
```

```
threshold = numeric()
```

```
for(i in 1:n)
```

```
{
```

```
  nn = neuralnet(Sin ~ Var, data = train, startweights = winit, hidden = 10, threshold = i/1000)
```

```
  pred_train = compute(nn, as.matrix(train$Var))$net.result
```

```
  pred_val = compute(nn, as.matrix(validation$Var))$net.result
```

```
  threshold[i] = i/1000
```

```
  print(i)
```

```
  mse_val[i] = mse(pred_val, validation$Sin)
```

```
  mse_train[i] = mse(pred_train, train$Sin)
```

```
}
```

```
plot(threshold, mse_val, type="o", ylab="MSE", xlab="Threshold ( $10^{-3}$ )", main = "Validation dataset")
```

```
plot(threshold, mse_train, type="o", ylab="MSE", xlab="Threshold ( $10^{-3}$ )", main = "Training dataset")
```

```
# Train network with full dataset
```

```
nn = neuralnet(Sin ~ Var, data = data, startweights = winit, hidden = c(10), threshold = 4/1000)
```

```
plot(nn)
```

```
# Predictions
```

```
plot(prediction(nn)$rep1, col="black", main = "Prediction (black) vs Data (red)")  
points(data, col = "red")
```

```
....
```